

```
In [1]: import csv
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report
from matplotlib.colors import ListedColormap
from sklearn import neighbors, datasets
```

```
/Users/shenstella/opt/anaconda3/lib/python3.9/site-packages/scipy/__init__.py:146: UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version 1.26.0)
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

## Load the data

```
In [2]: data = pd.read_csv('data.csv', parse_dates=['Intake-DateTime', 'Outcome-DateTime'])
data.dtypes
```

```
Out[2]: AnimalID          object
ID              object
Breed           object
Color           object
Gender          object
Name            object
Intake-DateTime    datetime64[ns]
Intake-Type        object
Intake-Condition   object
Intake-Age(days)    int64
Date-Of-Birth      object
Outcome-DateTime    datetime64[ns]
Outcome-Type        object
Outcome-Subtype     object
Outcome-Age(days)   int64
Age               object
IsAdopted          bool
Category           object
Intelligence-Ranking int64
Intelligence-Category object
Size-Category       object
Longevity          float64
Total-Cost($)       int64
Purchase-Cost($)    int64
Food-Cost($)        int64
Cost-Category       object
dtype: object
```

```
In [3]: data['ReFoundByShelter'] = data['ID'].str.contains('\+')
data
```

Out [3]:

	AnimalID	ID	Breed	Color	Gender	Name	Intake- DateTime	Intake- Type	Intake- Condition	Intake- Age(days)
0	A006100	A006100	Spinone Italiano	Yellow	Male	Scamp	2014-03-07 14:26:00	Public Assist	Normal	2190
1	A006100	A006100+	Spinone Italiano	Yellow	Male	Scamp	2014-12-19 10:21:00	Public Assist	Normal	2555
2	A006100	A006100++	Spinone Italiano	Yellow	Male	Scamp	2017-12-07 14:07:00	Stray	Normal	3650
3	A134067	A134067	Shetland Sheepdog	Brown	Male	Bandit	2013-11-16 09:02:00	Public Assist	Injured	12190
4	A141142	A141142	Labrador Retriever	Black	Female	Bettie	2013-11-16 14:46:00	Stray	Aged	11825
...	...	...	...	...	...	...	...	...	...	...
85791	A893431	A893431	Chihuahua	Tricolor	Female	Chili	2023-11-21 11:21:00	Public Assist	Normal	2920
85792	A893432	A893432	Chihuahua	Tan	Female	Coco	2023-11-21 11:21:00	Public Assist	Normal	2920
85793	A893452	A893452	Maltese	White	Female	Sophie	2023-11-21 13:38:00	Public Assist	Normal	2555
85794	A893529	A893529	Labrador Retriever	White	Female	Unknown	2023-11-22 14:26:00	Owner Surrender	Normal	30
85795	A893585	A893585	German Shepherd	Black	Male	Unknown	2023-11-23 20:19:00	Stray	Injured	30

85796 rows × 27 columns

In [4]: `y = data.iloc[:,26:27]`

Prepare the data, replace object data with numerical values

In [5]: `X = data.iloc[:,:].drop(columns=["IsAdopted", "ReFoundByShelter", "AnimalID", "ID", "Intake-Condition", "Intake-Age(days)"])`

Out[5]:

	Gender	Age	Intelligence-Ranking	Size-Category	Longevity	Total-Cost(\$)	Purchase-Cost(\$)	Food-Cost(\$)
0	Male	Adult	27	Large	9.00	18062	1725	5679
1	Male	Senior	27	Large	9.00	18062	1725	5679
2	Male	Senior	27	Large	9.00	18062	1725	5679
3	Male	Senior	6	Small	12.53	17469	465	3698
4	Female	Senior	7	Medium	12.04	18422	810	4819
...	...	...	...	...	...	...	...	...
85791	Female	Senior	67	Small	16.50	22640	588	4594
85792	Female	Senior	67	Small	16.50	22640	588	4594
85793	Female	Senior	59	Small	12.25	16073	650	2410
85794	Female	Baby	7	Medium	12.04	18422	810	4819
85795	Male	Baby	3	Large	9.73	15091	820	3895

85796 rows × 8 columns

```
In [6]: # Here we replace all the descriptive lable into numerical lable
gender = X.groupby('Gender').size()
X['Gender'].replace(['Female', 'Male'], [0,1], inplace=True)
SizeCategory = X.groupby('Size-Category').size()
X['Size-Category'].replace(['Large', 'Medium', 'Small'], [3,2,1], inplace=True)
Age = X.groupby('Age').size()
X['Age'].replace(['Baby', 'Young', 'Adult', 'Senior'], [1,2,3,4], inplace=True)
X
```

Out[6]:

	Gender	Age	Intelligence-Ranking	Size-Category	Longevity	Total-Cost(\$)	Purchase-Cost(\$)	Food-Cost(\$)
0	1	3	27	3	9.00	18062	1725	5679
1	1	4	27	3	9.00	18062	1725	5679
2	1	4	27	3	9.00	18062	1725	5679
3	1	4	6	1	12.53	17469	465	3698
4	0	4	7	2	12.04	18422	810	4819
...	...	...	...	...	...	...	...	...
85791	0	4	67	1	16.50	22640	588	4594
85792	0	4	67	1	16.50	22640	588	4594
85793	0	4	59	1	12.25	16073	650	2410
85794	0	1	7	2	12.04	18422	810	4819
85795	1	1	3	3	9.73	15091	820	3895

85796 rows × 8 columns

```
In [7]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [8]: # kNN is sensitive to the scale of features, so it's a good idea to normalize the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [9]: # Choose the number of neighbors (you can experiment with different values)
n_neighbors = 3

# Create kNN classifier
knn = KNeighborsClassifier(n_neighbors=n_neighbors)

# Train the model
knn.fit(X_train_scaled, y_train.values.ravel())

# Make predictions
y_pred = knn.predict(X_test_scaled)
```

```
In [10]: accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy}")
print("Classification Report:\n", report)
```

Accuracy: 0.8052447552447553  
Classification Report:

	precision	recall	f1-score	support
False	0.85	0.94	0.89	14543
True	0.17	0.07	0.10	2617
accuracy			0.81	17160
macro avg	0.51	0.50	0.50	17160
weighted avg	0.74	0.81	0.77	17160

## Using Cross Validation to Get the Best Value of k

```
In [11]: from sklearn.model_selection import cross_val_score
k_values = [i for i in range (1,31)]
scores = []

# Convert y_train to a NumPy array and then use ravel()
y_train_np = y_train.values.ravel()

X = scaler.fit_transform(X)


# Create kNN classifier
knn = KNeighborsClassifier(n_neighbors=n_neighbors)

# Train the model
knn.fit(X_train_scaled, y_train_np)
# Use cross_val_score
scores = []
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    score = cross_val_score(knn, X_train_scaled, y_train_np, cv=5)
    scores.append(np.mean(score))
```

We can plot the results with the following code

```
In [12]: import seaborn as sns
sns.lineplot(x = k_values, y = scores, marker = 'o')
plt.xlabel("K Values")
plt.ylabel("Accuracy Score")
```

Out[12]: Text(0, 0.5, 'Accuracy Score')



Since we know what could be the best K value, we update the model with the best K

```
In [14]: best_index = np.argmax(scores)
best_k = k_values[best_index]

knn = KNeighborsClassifier(n_neighbors=best_k)
knn.fit(X_train_scaled, y_train_np)
```

Out[14]: KNeighborsClassifier(n\_neighbors=26)

```
In [15]: y_pred = knn.predict(X_test_scaled)

accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy}")
print("Classification Report:\n", report)
```

Accuracy: 0.8476689976689976

Classification Report:

	precision	recall	f1-score	support
False	0.85	1.00	0.92	14543
True	1.00	0.00	0.00	2617
accuracy			0.85	17160
macro avg	0.92	0.50	0.46	17160
weighted avg	0.87	0.85	0.78	17160

evaluate with accuracy, precision, and recall

```
In [16]: from sklearn.metrics import precision_score, recall_score

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)

print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
```

Accuracy: 0.8476689976689976  
Precision: 1.0  
Recall: 0.0011463507833397019

```
In [17]: y_pred = knn.predict(X_test_scaled)
```

Use each feature alone - to fit a KNN model on the training set.

```
In [18]: import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier

# Feature names
names = ["Gender", "Age", "Intelligence-Ranking", "Size-Category", "Longevity", "Total-C

plt.figure(figsize=(15, 10))

# Iterate over each feature
for i in range(len(names)):
    plt.subplot(3, 3, i + 1)

    # Extract the current feature
    x = X[:, i].reshape(-1, 1)

    # Split the data
    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_stat

    # Convert y_train to a NumPy array and then use ravel()
    y_train_np = y_train.values.ravel()

    # Create KNN classifier
    knn = KNeighborsClassifier(n_neighbors=3)

    # Train the model
    knn.fit(x_train, y_train_np)

    y_pred = knn.predict(x_test)

    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)

    print("Evaluation for the feature: ", names[i])
    print("Accuracy:", accuracy)
    print("Precision:", precision)
    print("Recall:", recall)

    # Visualize decision boundary
    plt.scatter(x_test, y_test, label='Actual')
    plt.scatter(x_test, knn.predict(x_test), label='Predicted', marker='o', alpha=0.5)

    plt.title(f"Classification Relationship: {names[i]}")
    plt.xlabel(names[i])
    plt.ylabel('IsReFound')
    plt.legend()

plt.tight_layout()
plt.show()
```

/Users/shenstella/opt/anaconda3/lib/python3.9/site-packages/sklearn/metrics/\_classification.py:1318: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero\_division` parameter to control this behavior.

Evaluation for the feature: Gender

Accuracy: 0.8474941724941725

Precision: 0.0

Recall: 0.0

Evaluation for the feature: Age

Accuracy: 0.4244172494172494

Precision: 0.18489583333333334

Recall: 0.8139090561711884

Evaluation for the feature: Intelligence-Ranking

Accuracy: 0.8251165501165502

Precision: 0.17235494880546076

Recall: 0.03859380970576996

```
/Users/shenstella/opt/anaconda3/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

Evaluation for the feature: Size-Category

Accuracy: 0.8474941724941725

Precision: 0.0

Recall: 0.0

Evaluation for the feature: Longevity

Accuracy: 0.8374125874125874

Precision: 0.1685823754789272

Recall: 0.016813144822315627

Evaluation for the feature: Total-Cost(\$)

Accuracy: 0.8136363636363636

Precision: 0.18526543878656554

Recall: 0.06534199465036301

Evaluation for the feature: Purchase-Cost(\$)

Accuracy: 0.8057692307692308

Precision: 0.17395264116575593

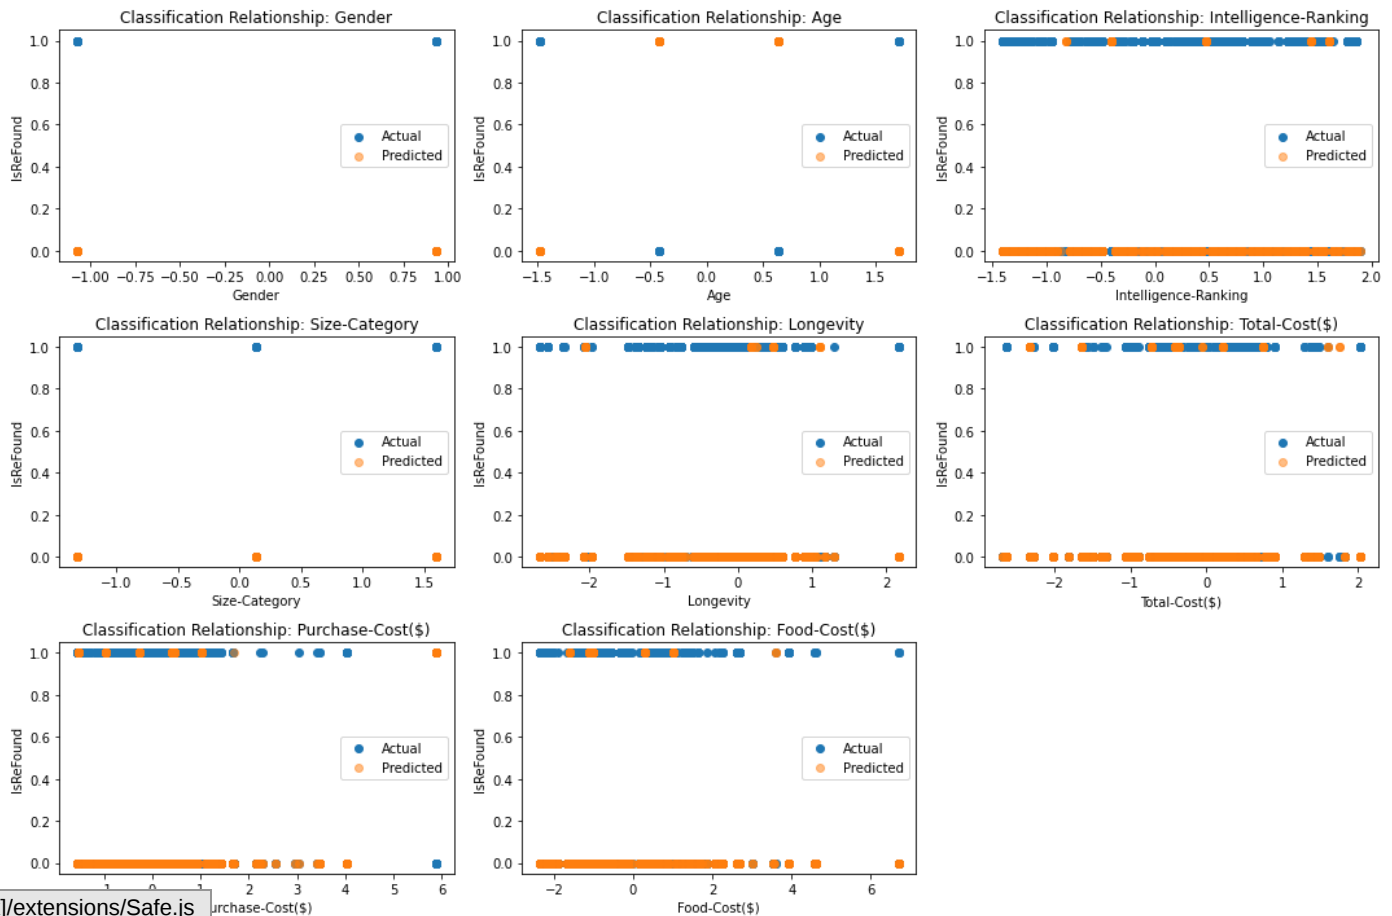
Recall: 0.0729843320596102

Evaluation for the feature: Food-Cost(\$)

Accuracy: 0.7314685314685314

Precision: 0.12448132780082988

Recall: 0.12609858616736722




# Visualization

```
In [19]: import matplotlib.pyplot as plt
from sklearn.decomposition import PCA

# Apply PCA to reduce the data to 2 dimensions (you can adjust this based on the number
pca = PCA(n_components=2)
X_train_pca = pca.fit_transform(X_train_scaled)

# Plot the data points in 2D space with color-coded classes
plt.scatter(X_train_pca[:, 0], X_train_pca[:, 1], c=y_train_np, cmap='viridis')
plt.title('PCA Visualization of Training Data')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```



## Conclusion & Discussion

From the Nearest Neighbor (kNN) Classifier model, we can see:

1. Age, Purchase cost, Food cost, Total-Cost do not affect if the pet will be refunded by shelter.
2. Size, Gender, Longevity and Intelligence level have a high influence on whether this dog will be refunded by shelter.
3. The model performs much better with all features rather than any individual ones. This quite makes sense since the model is better trained for classification.
4. The PCA plot suggests that to be refunded by shelter is a complex decision that can't be easily predicted by one or two features alone.

In [ ]: